



Stochastic Residual SAC: reward-based learning curves (alpha = 0.10)

Individual training seeds: faint raw returns, solid moving averages



Across-training-seed mean; interval is for raw episodic return



Recorded reward_return = sum of actual training rewards; entropy bonus and audit-only RPI penalties excluded.
Vertical dash: 5000-step warmup boundary. Random training paths; no new training or final test.
CI uses 3 training seeds, not pooled steps; descriptive only, not a simultaneous band or fixed-evaluation CI.